Jiahe Jin

zizi0123@sjtu.edu.cn google scholar homepage

Education

Shanghai Jiao Tong University, B.Eng. in Computer Science — Shanghai, China — 2022.09 – now

- **ACM Honors Class**, an elite research-oriented CS program.
- Zhiyuan Honors Program (top 5% of undergraduates).
- GPA: 3.93/4.3

Research Experience

Carnegie Mellon University, Research Intern, Advised by Prof. Chenyan Xiong — Pittsburgh, USA — 2025.04 – Now

Topics: Agentic Search; Deep Research Agents; Reasoning; Reinforcement Learning

Beneficial Reasoning Behaviors in Agentic Search and Effective Post-training to Obtain Them

- Build a pipeline to automatically identify reasoning behaviors in agentic search
- Propose a post-training method to prime models with reasoning behaviors to enable better performance in Reinforcement Learning.

Deepresearchgym: A Free, Transparent, and Reproducible Evaluation Sandbox for Deep Research

- Built a search agent sandbox with reproducible search API, and proposed a benchmark for deep research systems.

Deep Research Comparator: A Platform for Fine-grained Human Annotations of Deep Research Agents

- Developed a platform to host deep research agents and support side-by-side comparison and fine-grained human annotation for long reports evaluation.

Shanghai Jiao Tong University, Research Intern, Advised by Prof. Pengfei Liu — Shanghai, China — 2024.05 – 2025.05

Topics: Computer Use Agents, Vision-Language Models, Trustworthy LLMs

PC Agent: While You Sleep, AI Works–A Cognitive Journey into Digital World

- Built an infrastructure for collecting human-computer interaction trajectories, and a pipeline to train computer use agents from human demonstrations.

Efficient Agent Training for Computer Use

- Proposed a method to train computer use model by augmenting human trajectories with diverse action decisions from a frontier model, which outperforms using human data alone or direct distillation.

Behonest: Benchmarking honesty in large language models

- Introduced a benchmark assessing honesty in LLMs across awareness of knowledge boundaries, avoidance of deceit, and consistency in responses.

Publications

(indicates equal contribution)*

Beneficial Reasoning Behaviors in Agentic Search and Effective Post-training to Obtain Them
Jiahe Jin, *Abhijay Paladugu, Chenyan Xiong*
In submission to ACL 2026

Efficient Agent Training for Computer Use
Yanheng He, **Jiahe Jin***, Pengfei Liu*
In submission to ICLR 2026

Deep research comparator: A platform for fine-grained human annotations of deep research agents
Prahaladh Chandrahasan, **Jiahe Jin***, Zhihan Zhang*, Tevin Wang, Andy Tang, Lucy Mo, Morteza Ziyadi, Leonardo FR Ribeiro, Zimeng Qiu, Markus Dreyer, Akari Asai, Chenyan Xiong*
In submission to WWW 2026

Revisiting 3D LLM Benchmarks: Are We Really Testing 3D Capabilities?
Jiahe Jin*, *Yanheng He**, *Mingyan Yang**
Accepted by ACL 2025 Findings

PCAgent: While You Sleep, AI Works - A Cognitive Journey into Digital World
Yanheng He, **Jiahe Jin***, *Shijie Xia, Jiadi Su, Runze Fan, Haoyang Zou, Xiangkun Hu, Pengfei Liu*
Preprint

Deepresearchgym: A free, transparent, and reproducible evaluation sandbox for deep research
*João Coelho, Jingjie Ning, Jingyuan He, Kangrui Mao, Abhijay Paladugu, Pranav Setlur, **Jiahe Jin**, Jamie Callan, João Magalhães, Bruno Martins, Chenyan Xiong*
In submission to ICLR 2026

Behonest: Benchmarking honesty in large language models
*Steffi Chern, Zhulin Hu, Yuqing Yang, Ethan Chern, Yuan Guo, **Jiahe Jin**, Binjie Wang, Pengfei Liu*
Preprint

Generative ai act ii: Test time scaling drives cognition engineering
*Shijie Xia, Yiwei Qin, Xuefeng Li, Yan Ma, Run-Ze Fan, Steffi Chern, Haoyang Zou, Fan Zhou, Xiangkun Hu, **Jiahe Jin**, Yanheng He, Yixin Ye, Yixiu Liu, Pengfei Liu*
Preprint

Selected Course Projects

Revisiting 3D LLM Benchmarks: Are We Really Testing 3D Capabilities? — Computer Vision (A+)
- Identified an issue that some 3D LLM benchmarks could be easily solved by VLMs with rendered images, exposing ineffective evaluation the unique 3D capabilities.
- *first-author paper accepted by ACL 2025 Findings.*

Adaptive Length Control For Reasoning — Reinforcement Learning (A+)
- Applied a reward function that introduces token penalty according to question difficulty enable autonomous reasoning length control.

Teaching Experience

Data Structures (Honors), Teaching Assistant — Spring 2024

Awards

Ruiyuan-Sequoia Scholarship — 2025
- Awarded to top 0.5% of students in Zhiyuan Honor Program

Zhiyuan Honors Scholarship — 2023-2025
- Awarded to top 2% of students in SJTU

Shanghai Jiao Tong University Undergraduate Outstanding Scholarship (Class A) — 2025
- Awarded to students with top academic performance in SJTU

National High School Physics Competition (Zhejiang Division) — 2021
- First prize (84 students in Zhejiang Province)

Skills & Languages

Programming Languages: Python, Rust, C++, Java, Golang, Verilog.
Tools & Frameworks: verl, vLLM, LLaMA-Factory, Git, Docker, SLURM.